

Main

December 5, 2018

1 Program overview

- Look at relationships between income & different attributes.
- Extract important features from a Random Forest Model.

```
In [1]: import os
        os.chdir(r"C:\Users\Arnold\OneDrive\R_Python_working_directory")
        import pandas as pd
        import numpy as np
        from plotnine import *
        %matplotlib inline
        (theme_update(plot_background = element_rect(fill = "gold"),panel_background = element_
        colour = "blue",size = 1.99,linetype = "solid"),plot_title = element_text(hjust = 0.5))
```

2 Read in data

```
In [2]: df=pd.read_csv('cencus income.csv',na_values='?')
```

3 Data preview

```
In [3]: df.head()
```

```
Out[3]:
```

	age	workclass	fnlwgt	education	educational-num	marital-status	\
0	25	Private	226802	11th	7	Never-married	
1	38	Private	89814	HS-grad	9	Married-civ-spouse	
2	28	Local-gov	336951	Assoc-acdm	12	Married-civ-spouse	
3	44	Private	160323	Some-college	10	Married-civ-spouse	
4	18	NaN	103497	Some-college	10	Never-married	

	occupation	relationship	race	gender	capital-gain	capital-loss	\
0	Machine-op-inspct	Own-child	Black	Male	0	0	
1	Farming-fishing	Husband	White	Male	0	0	
2	Protective-serv	Husband	White	Male	0	0	
3	Machine-op-inspct	Husband	Black	Male	7688	0	
4	NaN	Own-child	White	Female	0	0	

	hours-per-week	native-country	income
0	40	United-States	<=50K
1	50	United-States	<=50K
2	40	United-States	>50K
3	40	United-States	>50K
4	30	United-States	<=50K

4 Data summary

```
In [4]: df.describe().T
```

```
Out [4]:
```

	count	mean	std	min	25%	75%	max
age	48842.0	38.643585	13.710510	17.0	28.0	45.0	99.0
fnlwgt	48842.0	189664.134597	105604.025423	12285.0	117550.5	162952.0	199999.0
educational-num	48842.0	10.078089	2.570973	1.0	9.0	12.0	16.0
capital-gain	48842.0	1079.067626	7452.019058	0.0	0.0	99999.0	99999.0
capital-loss	48842.0	87.502314	403.004552	0.0	0.0	4356.0	4356.0
hours-per-week	48842.0	40.422382	12.391444	1.0	40.0	45.0	99.0

5 Check the shape

```
In [5]: df.shape
```

```
Out [5]: (48842, 15)
```

6 Check % of rows with missing values

```
In [6]: np.mean(pd.isnull(df).any(axis=1))*100
```

```
Out [6]: 7.411653904426519
```

7% is low, so I'll drop the rows with missing values

7 Drop rows with missing values

```
In [7]: df.dropna(inplace=True)
```

8 Age binning

```
In [8]: df.age=pd.cut(df.age,bins=[16,20,30,40,50,60,70,100])
```

9 Remove columns that I'm not going to use

```
In [9]: df.drop(columns=['relationship','capital-gain','capital-loss','hours-per-week'],inplace=True)
```

9.1 Function to look at the relationships between different attributes & income


9.1.1 This function only works with categorical variables.

```
In [ ]: def relation(attr):
    tdf=df.pivot_table(values='fnlwgt',columns=['income'],index=attr,aggfunc=len,margins=True)
    tdf.iloc[:,-1,0]=tdf.iloc[:,-1,0].values/(tdf.All[:,-1]).values
    tdf.iloc[:,-1,1]=1-tdf.iloc[:,-1,0].values
    tdf.sort_values(by=['>50K'],inplace=True,ascending=False)
    tdf.iloc[1:,:2]=tdf.iloc[:,:2].applymap(lambda x: round(x*100))
    tdf=tdf.iloc[1:,:-1].stack().reset_index()
    tdf.rename(columns={0:'%'},inplace=True)
    return tdf
def rggplot(attr,size=8):
    graph=ggplot(tdf,aes(x=attr,y='%',fill=attr))+geom_bar(stat='identity',show_legend=False)
    size=size,format_string='{ }%')+coord_flip()+ggtitle('% of income > 50K')
    return graph
```

10 Gender vs Income

```
In [37]: tdf=relation('gender')
         tdf
```


```
C:\Anocanda\lib\site-packages\plotnine\layer.py:517: MatplotlibDeprecationWarning: isinstance(
return not cbook.iterable(value) and (cbook.is_numlike(value) or
```



Out[37]: <ggplot: (-9223371902897112309)>

```
In [13]: (ggplot(tdf,aes(x='income',y='%',fill='gender'))+geom_bar(stat='identity',position='dodge')+
  geom_text(aes(y=-5,label='gender'),position=position_dodge(width=0.9))+geom_text(aes(
  position_dodge(width=0.9),va='bottom', format_string='{}%'))
```

C:\Anocanda\lib\site-packages\plotnine\layer.py:517: MatplotlibDeprecationWarning: isinstance(
return not cbook.iterable(value) and (cbook.is_numlike(value) or



Out [13]: <ggplot: (133955488583)>

We can see that males are more likely to have high income than females.

11 Race VS Income

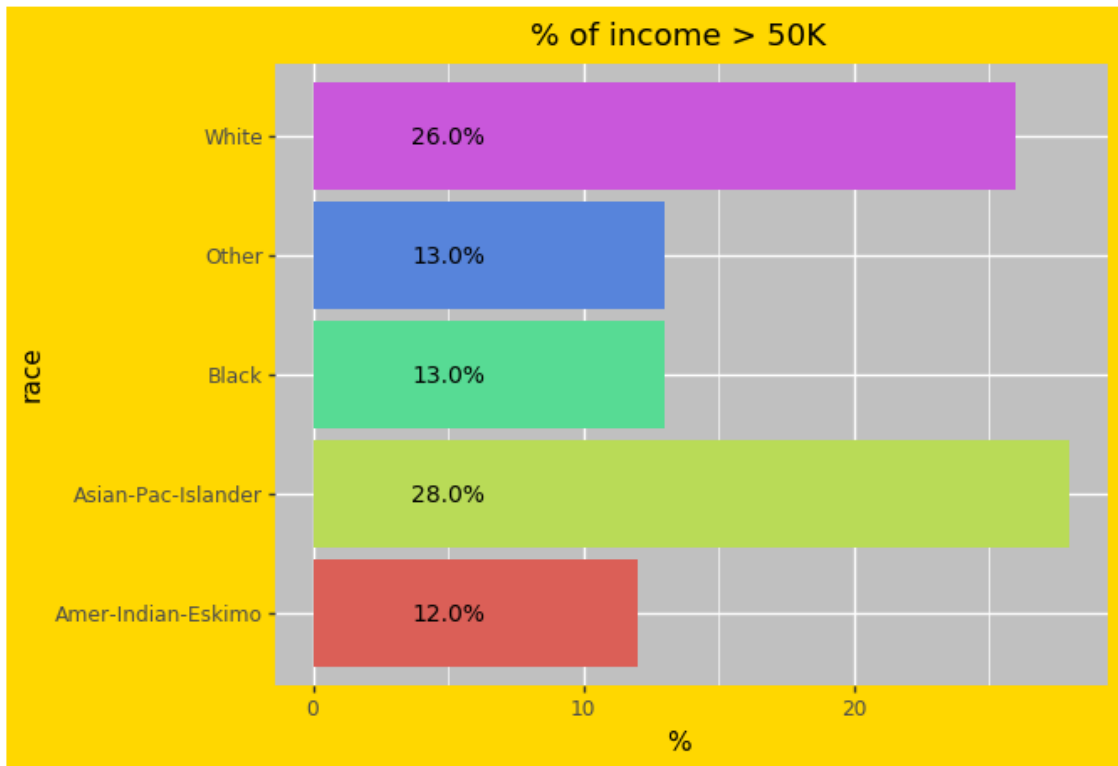
```
In [14]: tdf=relation('race')
tdf=tdf[tdf.income=='>50K']
tdf
```

```
Out [14]:
```

	race	income	%
1	Asian-Pac-Islander	>50K	28.0
3	White	>50K	26.0
5	Other	>50K	13.0
7	Black	>50K	13.0
9	Amer-Indian-Eskimo	>50K	12.0

```
In [15]: rggplot('race',10)
```

C:\Anocanda\lib\site-packages\plotnine\layer.py:517: MatplotlibDeprecationWarning: isinstance(
return not cbook.iterable(value) and (cbook.is_numlike(value) or



Out[15]: <ggplot: (133955726326)>

We can see that Asian-Pac-Islander & White races are more likely to have high income than other races.

12 Occupation VS Income

```
In [16]: tdf=relation('occupation')
         tdf=tdf[tdf.income=='>50K']
         tdf
```

```
Out[16]:
```

	occupation	income	%
1	Exec-managerial	>50K	48.0
3	Prof-specialty	>50K	45.0
5	Protective-serv	>50K	31.0
7	Tech-support	>50K	29.0
9	Armed-Forces	>50K	29.0
11	Sales	>50K	27.0
13	Craft-repair	>50K	23.0
15	Transport-moving	>50K	21.0
17	Adm-clerical	>50K	14.0
19	Machine-op-inspct	>50K	12.0
21	Farming-fishing	>50K	12.0

```

23 Handlers-cleaners >50K 7.0
25 Other-service >50K 4.0
27 Priv-house-serv >50K 1.0

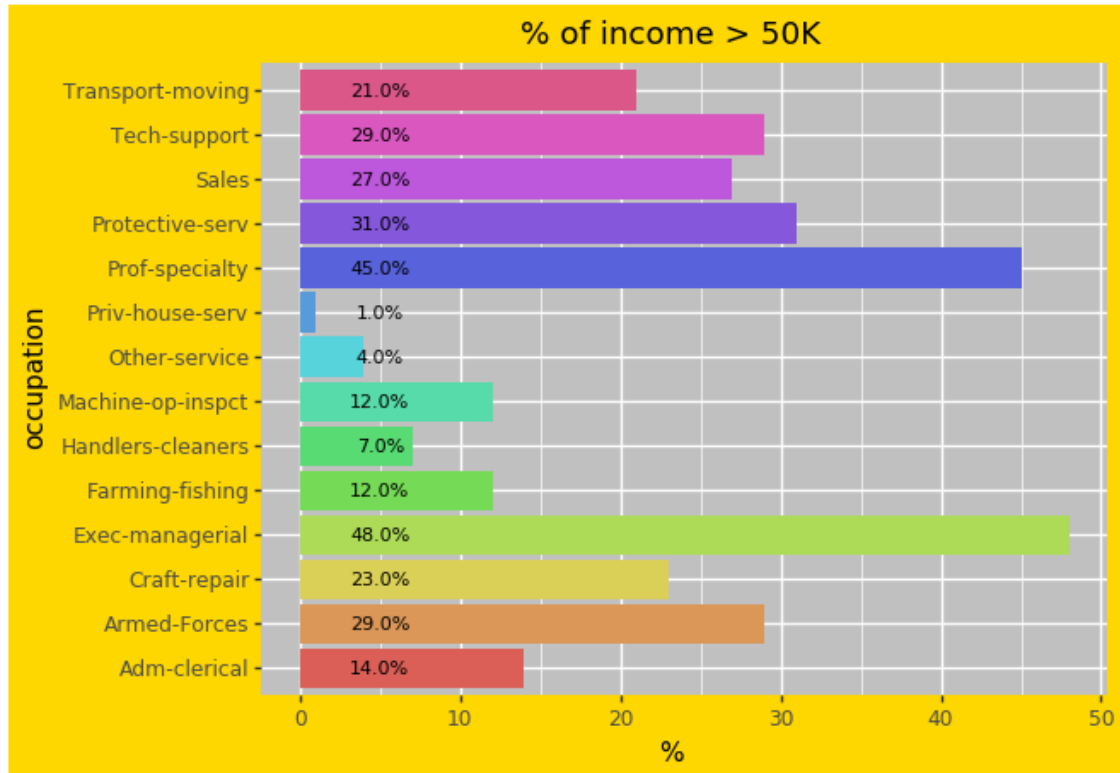
```

```
In [17]: rggplot('occupation')
```

```

C:\Anocanda\lib\site-packages\plotnine\layer.py:517: MatplotlibDeprecationWarning: isinstance(
return not cbook.iterable(value) and (cbook.is_numlike(value) or

```



```
Out[17]: <ggplot: (-9223371902899033207)>
```

We see that Exec-managerial, Prof-specialty, & Protective-serv are more likely to have high income than other occupations.

13 Workclass VS Income

```

In [18]: tdf=relation('workclass')
         tdf=tdf[tdf.income=='>50K']
         tdf

```

```

Out[18]:
      workclass income  %
1  Self-emp-inc  >50K 55.0

```

```

3      Federal-gov    >50K  39.0
5      Local-gov     >50K  30.0
7      Self-emp-not-inc >50K  28.0
9      State-gov     >50K  27.0
11     Private       >50K  22.0
13     Without-pay   >50K  10.0

```

```

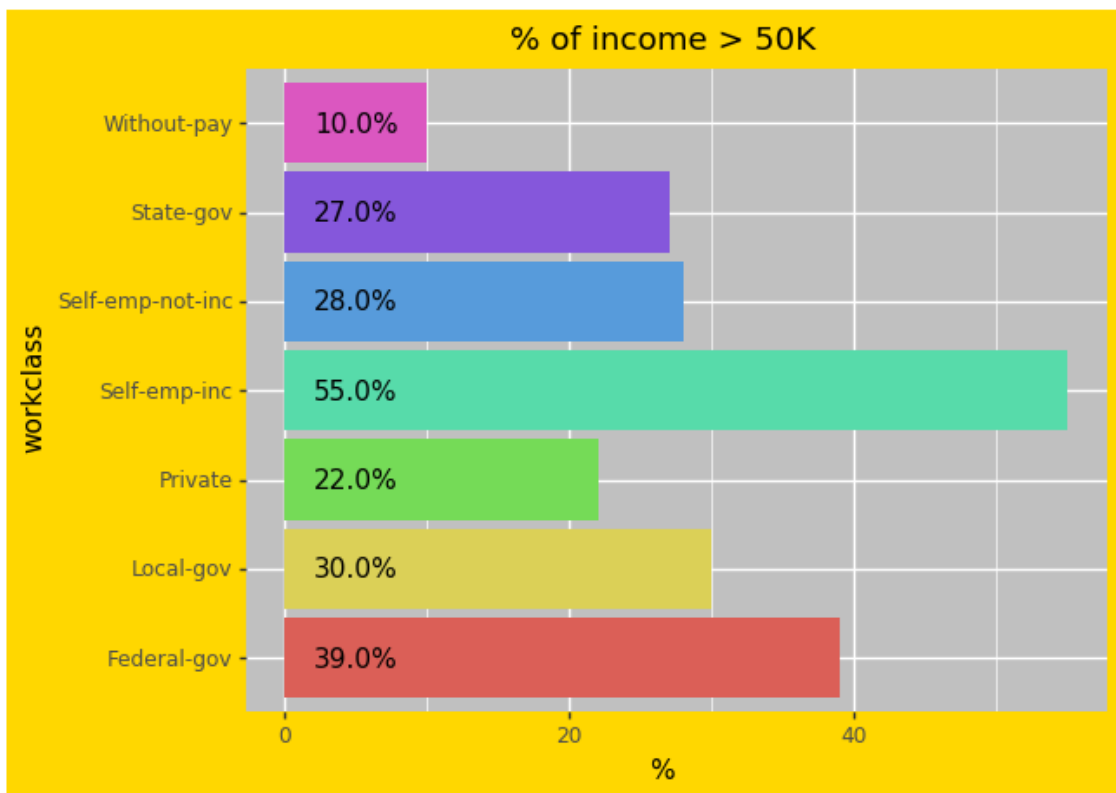
In [19]: (ggplot(tdf,aes('workclass','% ',fill='workclass'))+geom_bar(stat='identity',show_legend=
format_string='{ }%')+coord_flip()+ggtitle('% of income > 50K'))

```

```

C:\Anocanda\lib\site-packages\plotnine\layer.py:517: MatplotlibDeprecationWarning: isinstance(
return not cbook.iterable(value) and (cbook.is_numlike(value) or

```



```

Out[19]: <ggplot: (-9223371902898882921)>

```

We see that Self-emp-inc, Federal-gov, Local-gov are more likely to have high income than other workclass.

14 Marital-status VS Income

```

In [20]: tdf=relation('marital-status')
tdf=tdf[tdf.income=='>50K']
tdf

```




```
Out [20]:
```

	marital-status	income	%
1	Married-civ-spouse	>50K	45.0
3	Married-AF-spouse	>50K	44.0
5	Divorced	>50K	10.0
7	Married-spouse-absent	>50K	10.0
9	Widowed	>50K	9.0
11	Separated	>50K	7.0
13	Never-married	>50K	5.0

```
In [21]: rggplot('marital-status',10)
```

C:\Anocanda\lib\site-packages\plotnine\layer.py:517: MatplotlibDeprecationWarning: isinstance(
return not cbook.iterable(value) and (cbook.is_numlike(value) or



```
Out [21]: <ggplot: (133955472946)>
```

We see people who has the marital status of Married-civ-spouse, or Married-AF-spouse are much more likely to have high income than others.


15 Native-country VS Income

```
In [22]: tdf=relation('native-country')
         tdf=tdf[tdf.income=='>50K']
```

```
tdf
tdf=tdf.iloc[:20,:]
```

```
In [23]: rggplot('native-country')+ggtitle('% income > 50K top 20')
```

C:\Anocanda\lib\site-packages\plotnine\layer.py:517: MatplotlibDeprecationWarning: isinstance(
return not cbook.iterable(value) and (cbook.is_numlike(value) or



```
Out[23]: <ggplot: (133955817494)>
```

- We see that top 3 native country that have higher % of high income individuals are Taiwan, France, & India.
- U.S.A. as native country is ranked in the middle with 25% of people with high income.

16 Random Forest model

```
In [24]: from custom import tts
         from sklearn.ensemble import RandomForestClassifier as rfc
```

16.1 Getting data ready for modeling

```
In [25]: from sklearn.preprocessing import LabelEncoder as le
         from pandas import get_dummies as getd

In [26]: tdf=df.loc[:,['educational-num','workclass','marital-status','occupation','race','gender']]

In [27]: tdf.loc[:, 'age']=le().fit_transform(df.age)
         clm=['workclass','marital-status','occupation','race','gender','native-country','income']
         tdf=getd(tdf,columns=clm,prefix=clm,drop_first=1==1)
```

16.2 Split data into training and testing group

```
In [28]: xtr,xte,ytr,yte=tts(tdf,dep='income_>50K')

C:\Anocanda\lib\site-packages\sklearn\model_selection\_split.py:2069: FutureWarning: From vers
FutureWarning)
```

16.3 Model fitting

```
In [29]: mod=rfc(n_estimators=100).fit(xtr,ytr)
```

16.4 Look at feature importance

```
In [30]: imp=pd.DataFrame({'Attribute':tdf.columns[:-1],'Score':mod.feature_importances_.round(3)})
```

16.5 Clean up

```
In [31]: def clean(string):
         if '_' in string:
             return string[:string.find('_')]
         return string

In [32]: imp.Attribute=imp.Attribute.map(clean)

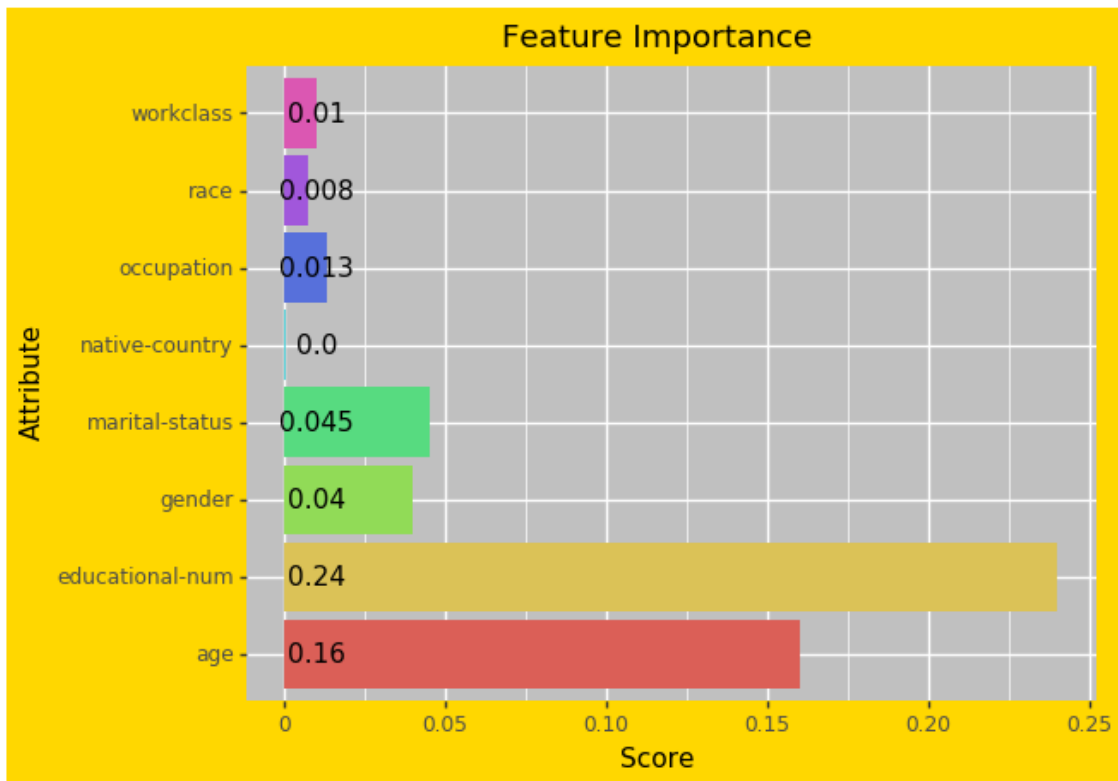
In [33]: tdf=imp.groupby('Attribute').mean().reset_index().sort_values('Score',ascending=1==0)
         tdf
```

```
Out[33]:
```

	Attribute	Score
1	educational-num	0.240000
0	age	0.160000
3	marital-status	0.045000
2	gender	0.040000
5	occupation	0.013077
7	workclass	0.010000
6	race	0.007500
4	native-country	0.000250

```
In [34]: ggplot(tdf,aes('Attribute','Score',fill='Attribute'))+geom_bar(stat='identity',show_l
         (label='round(Score,3)',y=.01))+coord_flip()+labs(title='Feature Importance',caption=
```

```
C:\Anocanda\lib\site-packages\plotnine\layer.py:517: MatplotlibDeprecationWarning: isinstance(
return not cbook.iterable(value) and (cbook.is_numlike(value) or
```



```
Out [34]: <ggplot: (133957643030)>
```

We see that education & age are by far the most important features according to this model.

16.6 Accuracy of the model

```
In [35]: mod.score(xte,yte)
```

```
Out [35]: 0.8231869103773585
```

```
In [36]: tdf.Attribute.unique()
```

```
Out [36]: array(['educational-num', 'age', 'marital-status', 'gender', 'occupation',
                'workclass', 'race', 'native-country'], dtype=object)
```